

Loading libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
from haversine import haversine
import seaborn as sns
```

```
In [2]: #pip install haversine
```

Loading Dataset & Performing EDA

```
In [3]: data_new=pd.read_csv("C:/Users/AYRUS/Desktop/nyc_taxi_trip_duration.csv")
```

```
In [4]: data_new.head()
```

Out[4]:

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	picku
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	1	-73.953918	
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	2	-73.988312	
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	2	-73.997314	
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	6	-73.961670	
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	1	-74.017120	

```
In [5]: data_new.columns
```

```
Out[5]: Index(['id', 'vendor_id', 'pickup_datetime', 'dropoff_datetime',
               'passenger_count', 'pickup_longitude', 'pickup_latitude',
               'dropoff_longitude', 'dropoff_latitude', 'store_and_fwd_flag',
               'trip_duration'],
              dtype='object')
```

```
In [6]: data_new.dtypes
```

```
Out[6]: id                object
vendor_id                int64
pickup_datetime          object
dropoff_datetime         object
passenger_count          int64
pickup_longitude         float64
pickup_latitude          float64
dropoff_longitude        float64
dropoff_latitude         float64
store_and_fwd_flag       object
trip_duration            int64
dtype: object
```

```
In [7]: data_new.isnull().sum()
```

```
Out[7]: id                0
vendor_id                0
pickup_datetime          0
dropoff_datetime         0
passenger_count          0
pickup_longitude         0
pickup_latitude          0
dropoff_longitude        0
dropoff_latitude         0
store_and_fwd_flag       0
trip_duration            0
dtype: int64
```

```
In [8]: data_new.nunique()
```

```
Out[8]: id                729322
vendor_id                2
pickup_datetime          709359
dropoff_datetime         709308
passenger_count          9
pickup_longitude         19729
pickup_latitude          39776
dropoff_longitude        27892
dropoff_latitude         53579
store_and_fwd_flag       2
trip_duration            6296
dtype: int64
```

```
In [9]: data_new.passenger_count.value_counts()
```

```
Out[9]: 1    517415
        2    105097
        5     38926
        3     29692
        6     24107
        4     14050
        0         33
        7          1
        9          1
        Name: passenger_count, dtype: int64
```

```
In [10]: data=data_new[data_new['passenger_count']!=0]
        data=data_new[data_new['passenger_count']<=6]
```

```
In [11]: data.passenger_count.value_counts()
```

```
Out[11]: 1    517415
         2    105097
         5     38926
         3     29692
         6     24107
         4     14050
         0         33
        Name: passenger_count, dtype: int64
```

```
In [12]: data['trip_duration'].sort_values(ascending=False)
```

```
Out[12]: 21813      1939736
         259437      86391
         119185      86387
         177225      86378
         496391      86377
         ...
        102646          1
        672240          1
        398887          1
        95823          1
        312992          1
        Name: trip_duration, Length: 729320, dtype: int64
```

```
In [13]: data.drop(data[data['trip_duration']==1939736].index,inplace=True)
```

In [14]: `data.describe()`

Out[14]:

	vendor_id	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	729319.000000	729319.000000	729319.000000	729319.000000	729319.000000	729319.000000
mean	1.535402	1.662038	-73.973513	40.750919	-73.973422	40.750919
std	0.498745	1.312405	0.069753	0.033594	0.069587	0.033594
min	1.000000	0.000000	-121.933342	34.712234	-121.933304	34.712234
25%	1.000000	1.000000	-73.991859	40.737335	-73.991318	40.737335
50%	2.000000	1.000000	-73.981758	40.754070	-73.979759	40.754070
75%	2.000000	2.000000	-73.967361	40.768314	-73.963036	40.768314
max	2.000000	6.000000	-65.897385	51.881084	-65.897385	51.881084

In [15]: `data.shape`

Out[15]: (729319, 11)

In [16]: *# Converting vendor_id and store_and_fwd_flag as category*
`data['vendor_id']=data['vendor_id'].astype('category')`
`data['store_and_fwd_flag']=data['store_and_fwd_flag'].astype('category')`

In [17]: *# Converting pickup_datetime and dropoff_datetime in datetime format*
`data['pickup_datetime']=pd.to_datetime(data['pickup_datetime'],format='%Y-%m-%d %H:%M:%S')`
`data['dropoff_datetime']=pd.to_datetime(data['dropoff_datetime'],format='%Y-%m-%d %H:%M:%S')`

In [18]: `data.dtypes`

Out[18]:

id	object
vendor_id	category
pickup_datetime	datetime64[ns]
dropoff_datetime	datetime64[ns]
passenger_count	int64
pickup_longitude	float64
pickup_latitude	float64
dropoff_longitude	float64
dropoff_latitude	float64
store_and_fwd_flag	category
trip_duration	int64
dtype:	object

In [19]: *# Creating new column and storing hour value*
`data['pickup_hour']=data['pickup_datetime'].dt.hour`

```
In [20]: # Creating new column and storing pickup time of day value
data['pick_time_of_the_day']=data.pickup_hour.apply(lambda y:(y%24+4)//4).replace
```

```
In [21]: # Creating new day_of_the_week column which stores the day of week
data['day_of_the_week']=data.pickup_datetime.apply(lambda x:x.day_name())
```

```
In [22]: # Converting trip duration in hours to second
data['trip_duration_in_hours']=data.trip_duration.apply(lambda x: x/3600)
```

```
In [23]: #Transforming day of week - Monday (0) to Sunday (6)
data['pickup_day_of_week']=data['pickup_datetime'].dt.dayofweek
```

```
In [24]: #Transforming pick up time of the day
condition=[data['pick_time_of_the_day']=='Late Night',
            data['pick_time_of_the_day']=='Early Morning',
            data['pick_time_of_the_day']=='Morning',
            data['pick_time_of_the_day']=='Afternoon',
            data['pick_time_of_the_day']=='Evening',
            data['pick_time_of_the_day']=='Night']

choice=[0,1,2,3,4,5]

data['pick_time_of_the_day_in_numbers']=np.select(condition,choice)
```

```
In [25]: # define a function to calculate distance
def calc_distance(df):
    pickup = (df['pickup_latitude'], df['pickup_longitude'])
    drop = (df['dropoff_latitude'], df['dropoff_longitude'])
    return haversine(pickup, drop)
```

```
In [26]: # creating a new column to store the distance value
data['total_distance'] = data.apply(lambda x: calc_distance(x), axis = 1)
```

```
In [27]: # calculating speed
data['speed'] = (data.total_distance/(data.trip_duration/3600))
```

In [28]: `data.head()`

Out[28]:

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	picku
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	1	-73.953918	
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	2	-73.988312	
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	2	-73.997314	
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	6	-73.961670	
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	1	-74.017120	

In [29]: `# Creating a dataframe to use for regression`
`data_features=['vendor_id', 'passenger_count', 'pickup_day_of_week', 'pick_time_of_t`

In [30]: `nyc_data=data[data_features]`
`nyc_data.head()`

Out[30]:

	vendor_id	passenger_count	pickup_day_of_week	pick_time_of_the_day_in_numbers	trip_duratio
0	2	1	0	4	
1	1	2	4	5	
2	2	2	6	4	
3	2	6	1	2	
4	1	1	2	1	

Segragating variables Independent and Dependent

In [31]: `# seprating independent and dependent variables`
`x=nyc_data.drop(['trip_duration_in_hours'],axis=1)`
`y=data['trip_duration_in_hours']`
`x.shape , y.shape`

Out[31]: ((729319, 6), (729319,))

Importing train test split

```
In [32]: # Importing Train test split
from sklearn.model_selection import train_test_split
train_x, test_x, train_y, test_y = train_test_split(x, y, random_state=56)
```

Importing Linear Regression & Mean Absolute Error

```
In [33]: from sklearn.linear_model import LinearRegression as LR
from sklearn.metrics import mean_absolute_error as mae
```

```
In [34]: # Creating instance of Linear Regression
lr = LR()
lr.fit(train_x, train_y)
```

Out[34]: LinearRegression()

```
In [35]: # Predicting over the Train Set and calculating error
train_predict = lr.predict(train_x)
k1 = mae(train_predict, train_y)
print('Training Mean Absolute Error', k1)
```

Training Mean Absolute Error 0.09736722896266335

```
In [36]: # Predicting over the Test Set and calculating error
test_predict = lr.predict(test_x)
k2 = mae(test_predict, test_y)
print('Test Mean Absolute Error', k2)
```

Test Mean Absolute Error 0.09902506089335923

```

In [37]: # creating data for bar plot
create_data={'Training_MAE':k1,
            'Test_MAE':k2}

display=list(create_data.keys())
value=list(create_data.values())

fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(display, value, color = 'blue',
        width = 0.4)

plt.xlabel("train and test error")
plt.ylabel("value of error")
plt.title("MAE of TEST AND TRAIN data")
plt.show

```

```

Out[37]: <function matplotlib.pyplot.show(close=None, block=None)>

```



```

In [38]: lr.coef_

```

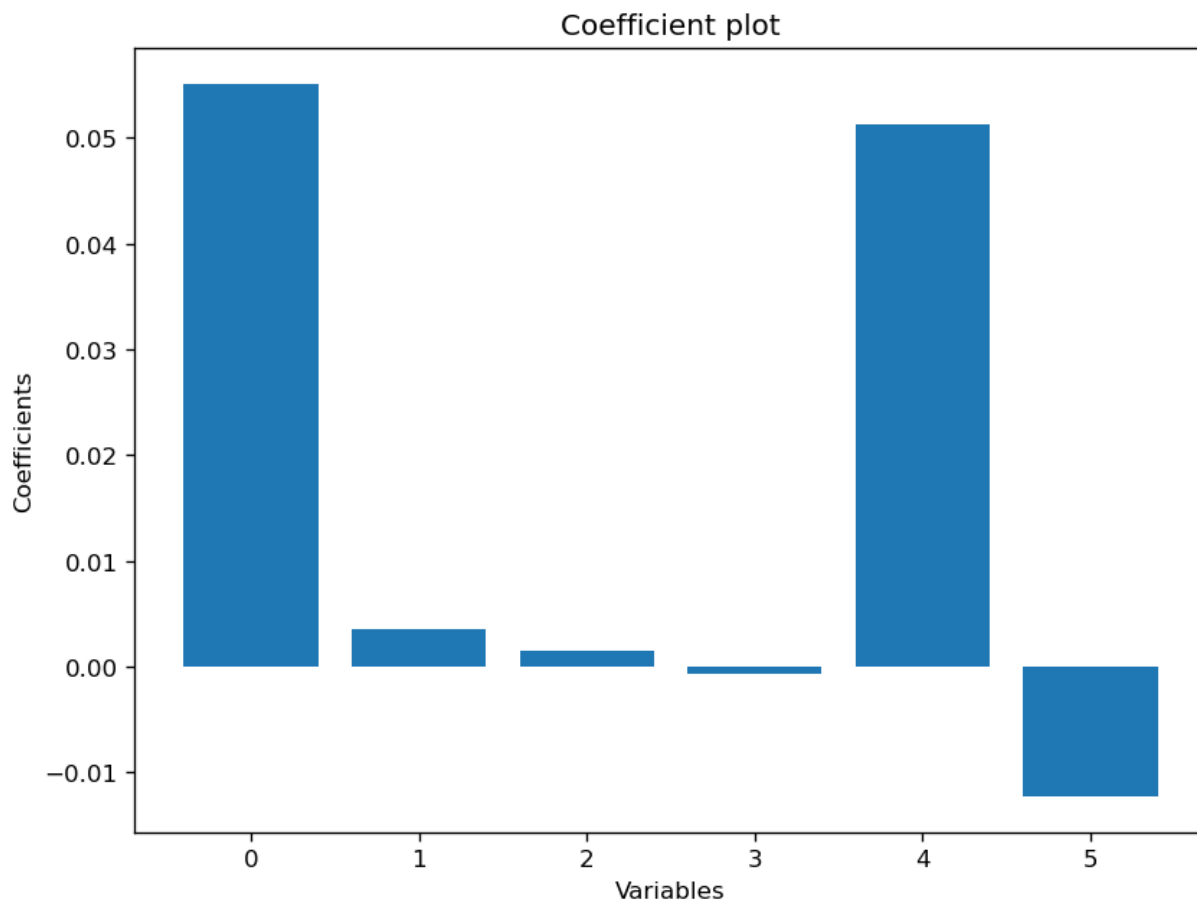
```

Out[38]: array([ 0.05510897,  0.00350898,  0.00146914, -0.00071647,  0.05121949,
                -0.01225439])

```



```
In [39]: plt.figure(figsize=(8, 6), dpi=120, facecolor='w', edgecolor='b')
x = range(len(train_x.columns))
y = lr.coef_
plt.bar( x, y )
plt.xlabel( "Variables")
plt.ylabel('Coefficients')
plt.title('Coefficient plot');
```



```
In [40]: # calculating the Residuals
residuals = pd.DataFrame({
    'fitted values' : test_y,
    'predicted values' : test_predict,
})

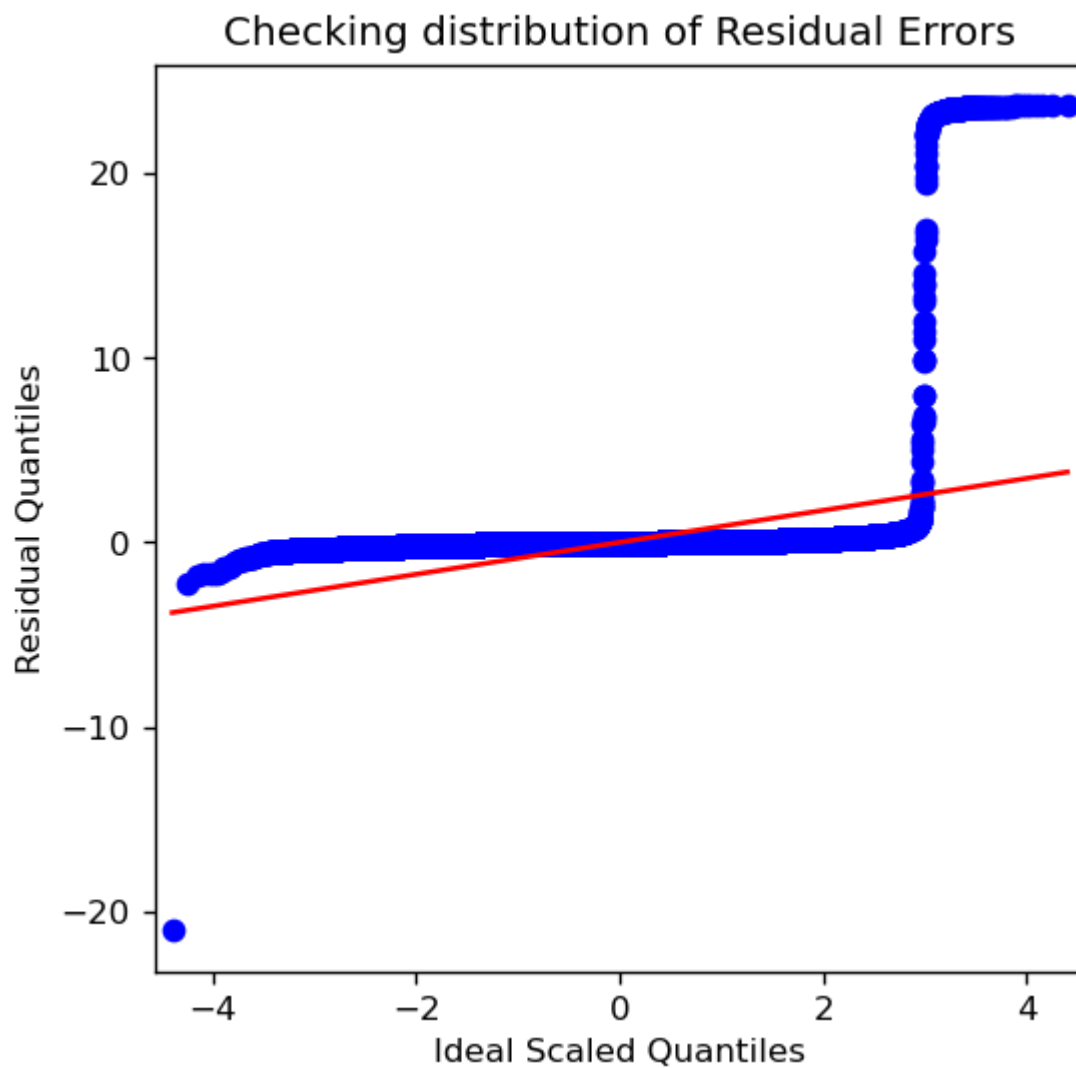
residuals['residuals'] = residuals['fitted values'] - residuals['predicted values']
residuals.head()
```

Out[40]:


	fitted values	predicted values	residuals
14908	0.668333	0.872062	-0.203728
223798	0.025278	0.177354	-0.152076
582522	0.049167	0.034962	0.014204
382010	0.133333	0.243796	-0.110463
619443	0.209722	0.222530	-0.012807

```
In [41]: # Importing qqplot
from statsmodels.graphics.gofplots import qqplot

## Plotting the QQ plot
fig, ax = plt.subplots(figsize=(5,5) , dpi = 120)
qqplot(residuals.residuals, line = 's' , ax = ax)
plt.ylabel('Residual Quantiles')
plt.xlabel('Ideal Scaled Quantiles')
plt.title('Checking distribution of Residual Errors')
plt.show()
```



```
In [42]: plt.figure(figsize=(10, 6), dpi=120, facecolor='w', edgecolor='b')
plt.hist(residuals.residuals, bins = 150)
plt.xlabel('Error')
plt.ylabel('Frequency')
plt.title('Distribution of Error Terms')
plt.show()
```



Importing ridge

```
In [43]: from sklearn.linear_model import Ridge
m = Ridge(alpha=0.01)
m.fit(train_x, train_y)
Ridge_train_score = m.score(train_x, train_y)
```

```
In [44]: # getting train score  
Ridge_train_score
```

```
Out[44]: 0.04641383298644963
```

Type *Markdown* and LaTeX: α^2

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```